

StegoCrypt — mathematical description

1. Goals & assumptions

Goals. Hide a secret payload (bytes) inside a cover image or video such that:

- The hidden bytes are **confidential** (requires password to decrypt).
- The modifications are **visually improbable** and **statistically difficult** to detect by casual inspection.
- The embedding/extraction is **deterministic** given the correct password (for reproducibility).
- Optionally, add **error-correction** for robustness to limited corruption.

Threat model.

- Passive adversary: can inspect the stego file, compare with known covers, or run automated steganalysis.
- Active adversary: may re-encode or compress (lossy) the file.
- StegoCrypt assumes sender and receiver can exchange the password securely out of band.

2. Payload packaging & cryptography

2.1 Payload format

Let the secret payload be a byte string M . We build an encapsulated payload P :

1. Salt: $s \leftarrow \text{Random}(\ell_s)$ (e.g., 16 bytes)
2. Key derivation: $K = \text{KDF}(\text{password}, s)$. StegoCrypt uses PBKDF2 (HMAC-SHA256) with iterations t and salt s .
3. AES-GCM encryption: choose random nonce n (e.g., 12 bytes), compute $C, T = \text{AES_GCM_Encrypt}(K, M; n)$, where T is the authentication tag.
4. Header: fixed magic bytes H and payload length $L = |C| + |T| + |n|$ encoded as fixed-size integer.

5. Optional ECC: if Reed–Solomon (RS) is enabled, apply RS encoding to the ciphertext+tag to produce E .

Final bytes to embed:

$$P = H \parallel L \parallel s \parallel n \parallel C \parallel T \text{ (or RS-coded version)}$$

Notes:

- AES-GCM provides confidentiality + integrity. On extraction, tag verification fails if either password wrong or ciphertext corrupted → “MAC check failed”.
- KDF parameters (salt size ℓ_s , iterations t) control brute-force resistance.

2.2 KDF formula

Use PBKDF2 with HMAC-SHA256:

$$K = \text{PBKDF2}_{\text{HMAC-SHA256}}(\text{password}, s, t, k)$$

where k is the desired key length (e.g., 256 bits). Recommended $t \geq 100,000$ (adjust for CPU).

3. Bit mapping, LSB embedding and capacity

3.1 Cover representation

- Images: each frame is an array of pixels with c color channels (typically $c = 3$ for RGB). A pixel channel is a byte $b \in \{0, \dots, 255\}$.
- Videos: sequence of F frames, each frame of resolution $W \times H$. Total pixel channels:

$$N_{\text{slots}} = F \cdot W \cdot H \cdot c$$

Each slot holds up to B bits if using B least significant bits (LSBs) per channel.

3.2 Capacity

Number of embeddable bits:

$$\text{Capacity_bits} = N_{\text{slots}} \cdot B$$

Number of embeddable bytes (payload capacity):

$$\text{Capacity_bytes} = \lfloor \frac{\text{Capacity_bits}}{8} \rfloor$$

Example. 320×240 video, 100 frames, $c = 3$, $B = 1$:

$$N_{\text{slots}} = 100 \cdot 320 \cdot 240 \cdot 3 = 23,040,000$$

Capacity bytes \approx 2.88 MB.

3.3 LSB embedding rule (per-slot)

For each payload bit $x \in \{0,1\}$ to embed into a channel byte b and using B LSBs:

- Let $b_{\text{msb}} = b \& \sim ((1 \ll B) - 1)$ (clear the low B bits).
- Let $b' = b_{\text{msb}} \mid x$ (if $B = 1$) or $b' = b_{\text{msb}} \mid$ next B payload bits.
- Write b' back.

This changes channel value by at most $2^B - 1$. With $B = 1$, $|b' - b| \in \{0,1\}$ — minimal perceptual difference.

4. Spreading (pseudo-random permutation) for stealth

4.1 Rationale

Sequentially writing payload bits in natural raster order produces visible/statistical artifacts (clusters, edges). Spreading pseudo-randomly distributes payload bits across slots to:

- Avoid contiguous clusters
- Reduce local statistical anomalies
- Make statistical detection harder

4.2 Deterministic permutation

Let $S = \{0, \dots, N_{\text{slots}} - 1\}$ be the ordered list of available slots (channel indices across frames). We compute a deterministic permutation π of S derived from the password and salt so both embedder and extractor reproduce it.

Common construction:

- Use a PRNG (cryptographic stream cipher or CSPRNG) seeded with $seed = \text{HMAC}(K_{\text{perm}}, s)$ where K_{perm} is derived from the password (or reuse K).
- Generate a random permutation via Fisher–Yates using the PRNG:
- $A = [0..N-1]$
- for $i = N-1$ down to 1:
- $j = \text{PRNG.next_integer}(0..i)$
- $\text{swap}(A[i], A[j])$
- Permutation $\pi(i) = A[i]$.

Because the header and salt are embedded sequentially at known low offsets (so the extractor can read salt/len early), the same permutation is reproducible.

4.3 Per-chunk permutation for streaming

For memory efficiency in video streaming:

- Process frames in chunks of C frames. For chunk k , define $N_k = C \cdot W \cdot H \cdot c$.
- Use seed derived from $seed_k = \text{HMAC}(K_{\text{perm}}, s \parallel k)$ so each chunk uses a deterministic permutation π_k .
- This allows embedding/extraction chunk-by-chunk without permuting the entire video at once.

5. Reed–Solomon ECC (optional)

5.1 Purpose & parameters

Reed–Solomon (RS) adds parity symbols to correct bit/byte corruptions. In StegoCrypt it operates on payload bytes after encryption but before embedding.

RS parameters: $\text{RS}(n, k)$ over $\text{GF}(2^8)$ where:

- message length k bytes,
- codeword length $n = k + 2t$,
- correctable symbol errors $t = \frac{n-k}{2}$.

If you set `rs_nsym = 32`, you add 32 parity bytes → correct up to 16 byte errors.

5.2 Trade-off

- ECC increases the number of bytes to embed by factor $\frac{n}{k}$.
- Larger ECC \rightarrow better robustness to corruption but larger stego payload \rightarrow more capacity required and potentially larger detectability footprint.

5.3 Integration

Pipeline:

$$M \xrightarrow{\text{encrypt}} C \xrightarrow{\text{RS encode}} E \xrightarrow{\text{embed}} \text{stego}$$

On retrieval: extract bits \rightarrow bytes \rightarrow RS decode \rightarrow AES-GCM decrypt \rightarrow verify tag.

If RS decoding fails or AES-GCM tag fails, extraction aborts (integrity failure).

6. Video-specific considerations (streaming, codecs)

6.1 Streaming chunking

Frames are processed in chunks C . For chunk k the mapping of global slot index to (frame, y, x, channel) is:

$$\begin{aligned} \text{slot_index} \mapsto (f = \lfloor \frac{\text{slot_index}}{W \cdot H \cdot c} \rfloor, y = \lfloor (\text{slot_index} \bmod (W \cdot H \cdot c)) / (W \cdot c) \rfloor, x \\ = \lfloor ((\text{slot_index} \bmod (W \cdot H \cdot c)) \bmod (W \cdot c)) / c \rfloor, ch = (\cdot)) \end{aligned}$$

When using per-chunk permutations, reuse the chunk-local seed.

6.2 Lossless vs lossy codecs

- Lossless codecs (FFV1, libx264rgb lossless mode) preserve decoded pixel values exactly \rightarrow ideal for LSB steganography.
- Lossy codecs (standard H.264 with chroma subsampling, motion compression) will change LSBs unpredictably \rightarrow will usually break simple LSB extraction unless heavy ECC and redundancy are used.

7. Steganalysis & detectability

7.1 Simple statistical tests

- **LSB plane analysis.** For natural images the LSBs behave close to random, but certain patterns differ by region. Embedding uniformly can increase entropy or break expected correlations.
- **Chi-square test** (or RS analysis) can detect changes to LSB distributions.
- **Sample pairs analysis** looks for deviations in parity frequency.

7.2 Effect of spreading

Permutation spreading reduces local cluster changes, making spatial statistics more natural; however global statistics (histograms, pairwise correlations) may still change.

7.3 Metrics

- Let $\Delta(b)$ denote distribution of bit flips across intensities or local neighborhoods. A strong stealth metric is minimizing the KL divergence between the empirical LSB distribution of cover and stego:

$$D_{\text{KL}}(P_{\text{cover}} \parallel P_{\text{stego}}) = \sum_x P_{\text{cover}}(x) \log \frac{P_{\text{cover}}(x)}{P_{\text{stego}}(x)}.$$

Spreading + $B=1$ minimizes detectable divergence vs larger B .

7.4 Practical recommendation

- Use $B = 1$ and spread (default) for best stealth.
- Avoid embedding near image edges or uniform large flat regions (but spreading already mitigates this).
- Do not re-encode with lossy codecs after embedding.

8. Error probability & reliability (back-of-envelope)

Assume an adversary or transmission causes independent per-channel bit-flip probability p . For $B = 1$ and payload of m bytes, expected number of corrupted bytes depends on bit-to-byte mapping and deinterleaving (spreading acts as interleaver).

If RS corrects up to t byte errors, probability of successful decode is:

$$P(\text{success}) = \sum_{e=0}^t \binom{m+n_{\text{par}}}{e} q^e (1-q)^{(m+n_{\text{par}})-e}$$

where q is the byte-error probability (derived from bit flip model). This is approximative.

9. Algorithms & pseudocode

9.1 Embed (high-level)

Input: cover (image/video frames), password, message bytes M, params: B, use_spread, use_rs, rs_nsym

1. P = build_payload(M, password, use_rs=use_rs, rs_nsym=rs_nsym)
2. header = H || len(P) || salt
3. bits = bytes_to_bitstream(header || P)
4. if use_spread:
 - compute permutation(s) pi (global or per-chunk) seeded by password+salt
 - slots = apply_permutation(pi) # an ordered list of channel indices to write into
- else:
 - slots = [0..N_slots-1]
5. for i, bit in enumerate(bits):
 - slot = slots[i]
 - write_bit_into_slot(slot, bit, B)
6. save stego cover (frames->container). Use lossless codec.

9.2 Extract (high-level)

Input: stego file, password, params B, use_spread, use_rs

1. Read first N_header_slots sequentially to recover header(H, len, salt)
2. derive K, permutation seed from password+salt
3. build slots via permutation (same as embed)

4. read bitstream by reading low-B bits from slots
5. reconstruct bytes, optionally RS-decode
6. attempt AES-GCM decrypt with K and nonce; if tag valid, output M

10. Security analysis

10.1 Confidentiality

- AES-GCM with strong KDF (PBKDF2 with sufficient iterations) gives standard IND-CPA/AE guarantees: without password, ciphertext indistinguishable from random (subject to nonce reuse rules).
- Salt and nonce must be random and unique per embedding.

10.2 Integrity

- AES-GCM tag prevents undetected modification. If tag verification fails extraction reports failure; RS cannot correct arbitrary malicious tampering beyond its correction capability.

10.3 Stealth

- Confidentiality does not imply stealth: ciphertext has high entropy and if embedded densely may increase local entropy and be detectable.
- Spreading reduces local detectability; limiting per-channel modification to 1 LSB is the primary practical stealth tradeoff.

10.4 Brute-force & password strength

- Attacker can attempt password guesses. Each attempt requires full extraction and AES verification. Use high KDF work factor and encourage strong passwords.

11. Parameter recommendations (practical)

- LSB $B = 1$ (default) — best stealth.
- Use spread = True.
- RS only if you anticipate lossy transforms. If used: rs_nsym set based on expected corruption; 32 parity bytes correct up to 16 byte-errors.
- KDF: PBKDF2-SHA256 with iterations $t = 200,000$ (adjust to CPU), salt 16 bytes.

- AES-GCM: 256-bit key, 12-byte nonce, 16-byte tag.
- Chunk size for video: $C = 30\text{--}90$ frames, depending on memory.

12. Example numeric scenario

320×240, 60 frames, $c = 3$, $B = 1$:

- slots: $60 \cdot 320 \cdot 240 \cdot 3 = 13,824,000\text{bits} \rightarrow \text{capacity} \approx 1.73 \text{ MB payload bytes.}$
- After AES-GCM + header + salt (assume 128 bytes overhead), usable payload $\approx 1.73 \text{ MB} - 128 \text{ B.}$